

AI Integration and Economic Effects

Research Note 1 — First execution of the canonical pipeline: retrieval audit, validated adoption panel, AI Integration Surprise, and diffusion dynamics · 25 August 2026 · Evidence grade: EXPLORATORY

Summary. This note reports the first end-to-end execution of the research framework frozen in the design checklist: source retrieval with manifests and fingerprints, validation against published values, canonical panel construction, and a pre-declared exploratory analysis family with permutation and leave-one-country-out (LOCO) diagnostics. Three results stand out. **(1)** The primary adoption source (Microsoft AI Diffusion, 147 economies, three periods) was retrieved first-party and passed all 15 validation checks, including exact reproduction of eight published values. **(2)** The first AI Integration Surprise (AIS) estimates show a steep, extremely robust income gradient in adoption (logit slope 0.44 per log-dollar of GDP per capita, $R^2=0.71$, $N=142$) — and the largest *negative* surprises concentrate almost entirely in provider-restricted economies (mean AIS -0.44 vs 0.02 elsewhere, permutation $p=0.007$), giving the first within-pipeline confirmation of the availability confounder (G3) the design anticipated. **(3)** The additional finding: between H1 2025 and Q1 2026, global AI diffusion **diverged in absolute terms** — each percentage point of initial adoption predicts 0.14 pp of additional growth over the window (search-adjusted permutation $p=0.0005$ across the four-spec dynamic family; LOCO range [0.130, 0.147]). The economic-outcome (EPS) regression — the design's priority-1 test — could **not** be run: every outcome source is unreachable from this execution environment, an audit result reported in Section 2.

1. What was executed

The pipeline implements the layered architecture of the research skeleton: *raw* (immutable downloads), *manifest* (URL, UTC timestamp, bytes, SHA-256, license, fingerprint), *staging* (deterministic transforms, Mac-Roman decoding of the source CSV, ISO3 concordance), *panel* (147 economies, unique keys, explicit missingness), *exploration* (pre-declared spec family, machine-readable search ledger), and *reports* (this note; no manually typed numeric results — every statistic is interpolated from the pipeline's JSON outputs). Deterministic seed 20260825 throughout; Microsoft repository vintage pinned at git HEAD 8fe84a3b797e.

2. Retrieval audit: what the registry could and could not deliver

Of the 17 registry sources, the primary adoption source was retrieved **first-party** (Microsoft's official ai-diffusion-report repository: country CSV plus all four report PDFs, fingerprints matching the registry exactly). Structural covariates were retrieved from Frictionless Data mirrors of the World Bank API (GDP current US\$ to 2023; population to 2024) with the mirror channel disclosed. Ten sources are unreachable from this environment: the network egress policy allows GitHub, microsoft.com and package registries only. Every blocked source is recorded in the manifest with status *blocked_by_egress* rather than silently substituted.

Status	Sources	Consequence
Retrieved, first-party	S01 (Q1'26 + H1'25 report PDFs), S02 (country CSV + README)	Primary adoption measure fully available
Retrieved, mirror	S15 GDP (Frictionless/World Bank), population, ISO3/Region, etc.	Structural covariates for AIS residual model
Blocked by egress	S03, S04, S05, S07, S08, S10, S11, S12, S13, S14	No Anthropic/OpenAI cross-provider check (H4 criterion untestable); no E

Consequences for the frozen design: the cross-provider consistency requirement for a Provisional grade cannot be met from this environment; and the priority-1 regression ($AIS \times \text{sector exposure} \rightarrow \text{productivity surprise}$) cannot be estimated at all, because ec.europa.eu (Eurostat nama_10_lp_a21, sts_sepr_m), imf.org (WEO April 2026 and the October 2024 forecast baseline) and ilo.org (white-collar denominators) are policy-blocked. This note therefore delivers the treatment-side measures and dynamics, and specifies the outcome branch for execution in an environment with statistical-agency access.

3. Validation: 15/15 checks passed

Following the retrieval-validation protocol (skeleton section 12.1): unique-key, range, denominator and missingness checks, plus reproduction of values visibly published in the Q1 2026 report — all passed. Reproduced exactly: UAE 70.1, Singapore 63.4, Norway 48.6, Ireland 48.4, United States 31.3, China 16.4 (Q1 2026 AI user share, percent), UAE 64.0 and Singapore 60.9 (H2 2025), and the United States' published rank of 21st. The population-weighted world mean lands at 17.60 against the published 17.8 — the deviation is expected and disclosed: the report weights


by population aged 15-64, while only total population is retrievable through the allowed channel. Five economies lack GDP per capita in the World Bank mirror and are dropped from residual models with explicit listing (Taiwan, Venezuela, Eritrea, South Sudan, French Guiana).

4. AI Integration Surprise: the design's core measure, first estimates

Per the frozen answers (D1/D5), AIS is the residual of adoption after structural predictors, estimated with a transparent linear model and a rank check. The frozen covariate list (D4) is only partially retrievable; the constrained model — logit adoption on log GDP per capita, log population and M49 region fixed effects — is a logged deviation, and AIS is computed as the **jackknife** (leave-one-out) residual so no observation scores its own fit, honoring the cross-fitting requirement (skeleton 4.1).


Spec	Dependent	Coef. log GDP pc (HC1 SE)	R ²	Note
A1	level	6.56	0.655	levels, no controls
A2	logit	0.418 (0.022)	0.699	primary functional form
A3	logit	0.419	0.699	+ log population
A4 (primary)	logit	0.444 (0.040)	0.706	+ region FE; log pop coef 0.007 (0.029) — no scale effect
A5	rank	Spearman rho = 0.826	—	robustness

The income gradient is steep and stable: LOCO range for the A4 slope [0.422, 0.452] (extremes when dropping ARE and BDI). One log-unit of income (~2.7x) moves adoption by ~0.44 logit — for a country at 10 percent adoption, roughly seven additional points. Population is irrelevant once income and region are held (coefficient 0.007, SE 0.029): consistent with the source's population normalization doing its job, and with adoption being a per-capita structural trait, not a country-scale artifact (the skeleton's 'country scale' confounder row behaves as designed).



The AIS league table. Over-adopters given income, size and region: UAE (+1.67) and Singapore (+1.10) remain exceptional even after conditioning — their leads are not income artifacts. The rest of the positive tail is informative: Jordan (+0.93), Lebanon (+0.91), Vietnam (+0.77) and Colombia (+0.56) out-adopt their structural predictions — middle-income economies with young, digitally active workforces; France (+0.66), Spain (+0.65) and New Zealand (+0.63) lead among advanced economies.

Post-hoc finding C (flagged as post-hoc in the search ledger). The negative tail is not random: Turkmenistan (-1.27), Cuba (-1.17), Russia (-1.10), Belarus (-0.78), Armenia (-1.07), Ukraine (-0.58). Prompted by this pattern, a post-hoc group test was run for economies where the major Western AI providers are unavailable or restricted (AFG, BLR, CHN, CUB, IRN, RUS, SYR): mean AIS -0.44 versus 0.02 for all others — difference -0.46, permutation $p = 0.0075$; China alone sits at -0.40. This is the design's G3 confounder (language, availability, provider policy) surfacing empirically on the first pass: **provider-based adoption telemetry measures access-conditional adoption, and negative AIS values in restricted markets must never be read as low latent demand.** It vindicates the expected-answers' B3 position (handle restricted markets separately) over this pipeline's initial exclusion-only stance, and upgrades that decision's priority for the next design revision.




5. The additional finding: absolute divergence in global AI diffusion

With three reporting periods (H1 2025 → Q1 2026, roughly nine months of coverage), the pre-declared dynamic family asks whether diffusion converges or diverges across economies. All four specifications were declared before estimation and every run is in the search ledger; the headline claim is search-adjusted by max-statistic permutation across the family.

Spec	Model	Focal coefficient	Perm. p	R ²	LOCO range
B1 (headline)	$\Delta pp \sim \text{initial level}$	0.136 pp per pp	<1e-4	0.667	[0.130, 0.147]
B2	$\Delta pp \sim \log \text{ GDP pc}$	0.884 pp per log\$	<1e-4	0.561	[0.846, 0.901]
B3	$\Delta pp \sim \text{initial} + \log \text{ GDP pc}$	initial: 0.100	<1e-4	0.693	[0.081, 0.117]
B4	$\log \text{ growth} \sim \log \text{ initial}$	-0.0207	0.0019	0.067	[-0.0233, -0.0189]

Reading. An economy ten points higher in H1 2025 gained on average 1.4 additional points by Q1 2026 (B1); the effect survives an income control (B3: 0.100 conditional on log GDP per capita) and label permutation at the family level (max $|t| = 12.6$, search-adjusted $p = 0.0005$ over 4 specifications). B4 shows mild *proportional* convergence (laggards grow faster in percentage terms, coefficient -0.0207) — but far too weak to close absolute gaps: the percentage-point distance between leaders and laggards is widening, which is the economically relevant margin for any AIS-based treatment. This quantifies, and income-adjusts, what the source report states descriptively (a Global North-South gap of 9.8 → 10.6 → 12.1 points across the three periods, values we validated in Section 3).

Honest mechanical caveat. Under a logistic diffusion curve, absolute gains peak at 50 percent penetration; with most leaders still below 50 percent, absolute divergence is what standard S-curve diffusion predicts at this phase, not yet evidence of a permanent Matthew effect. The decisive question — whether laggards' takeoff is merely delayed (the historical norm for technology diffusion, cf. Comin-Hobijn) or structurally suppressed (the source report's own electricity/connectivity/skills reading) — is answerable within this design once two or three more report vintages exist. The design's variance/persistence statistics (skeleton 7.1) transfer directly, and the prospective 2027-2028 holdout applies to this finding as much as to the EPS branch.



6. The blocked branch: what the EPS regression needs

The priority-1 test remains: EPS on jackknifed AIS interacted with sector cognitive exposure, with country×time and sector×time fixed effects and wild-cluster bootstrap. To execute it, the environment needs: (i) ec.europa.eu — Eurostat nama_10_lp_a21 (real GVA per hour by NACE section, exposed J/M/K vs controls) and sts_sepr_m (monthly services output, intermediate outcome); (ii) imf.org — WEO April 2026 actuals plus the October 2024 pre-AI forecast vintage for growth surprises; (iii) ilo.org — employment by ISCO-08 for the white-collar denominator; (iv) huggingface.co and openai.com — Anthropic AEI and OpenAI Signals for the cross-provider requirement. The AIS measures produced here are the treatment-side input to that regression and are versioned in the panel for reuse. Timing note: with adoption first measured in H1 2025, outcome data through 2023 (the mirror vintage available here) precedes treatment measurement and was correctly not used — the no-look-ahead assertion in the validation layer would have rejected it.

7. Parallels with existing research

Microsoft AI Economy Institute (2026), the source reports. Everything here is consistent with, and validated against, their published aggregates. The additions: a formal income-residualized surprise measure (their reports publish raw shares and narrative income commentary, not a conditional model), the jackknife/cross-fit treatment, permutation and LOCO inference, and the provider-restriction group test.

The Jagged Global Economy (arXiv 2607.05404, 2026). The closest neighbor; it links national AI exposure to adoption across the same three provider ecosystems and reports steep, super-linear gradients (their headline: a 0.10 increase in national AI exposure associates with a ~12x increase in Claude usage per working-age capita). Our logit income slope (0.42-0.44 per log-dollar) is the Microsoft-telemetry analog of that steepness, on a broader 142-economy panel and with the availability confound explicitly separated. Their exposure measure would be the natural CognitiveExposure input when the EPS branch runs.

St. Louis Fed (2026) and Bick-Blandin-Deming-style surveys. These report an adoption gap (US ~43 percent of workers vs 26-36 percent in Europe, early 2026) and, at industry level, faster productivity growth in more-adopting industries (~3.2 pp cumulative US-Europe differential since 2022). Note the measurement difference our validation makes visible: telemetry-based population shares (US 31.3, France 47.8) rank countries differently than worker-survey shares — France out-adopts the US in Microsoft telemetry while surveys put the US first among workers. Reconciling telemetry vs survey instruments is exactly the provider-reconciliation problem the design's C2 answer anticipates, now with concrete numbers attached.

Technology-diffusion literature (Comin-Hobijn; ICT wave). The absolute-divergence-with-proportional-convergence pattern replicates the classic finding for past general-purpose technologies: adoption *lags* eventually close while intensity gaps persist. AI's diffusion is tracking the historical GPT template so far — which is itself a useful null against stronger 'AI is different' claims in either direction.

Productivity J-curve (Brynjolfsson-Rock-Syverson). Nothing here tests productivity; but the steep adoption gradients plus the J-curve literature jointly explain why the design's priors put mass on lagged, mechanism-close outcome effects rather than contemporaneous GDP effects.

8. Evidence grading and limitations

Grade: Exploratory (reproducible from immutable raw data; plausible mechanisms; no claim beyond hypothesis generation), per the project ladder. Binding limitations: single-provider adoption telemetry (cross-provider consistency untestable from this environment — the H4 promotion criterion is failed by construction, not by the data); GDP per capita is nominal 2023 from a mirror channel (the frozen design's covariate list D4 is only partially honored — a logged deviation); total rather than working-age population; three time periods only; and the dynamic findings concern the *treatment* variable's evolution, not economic outcomes. The divergence finding's S-curve caveat (Section 5) is part of the claim, not a footnote to it.

9. Provenance

Artifact	Bytes	SHA-256 (first 16)
S02 Microsoft AI Diffusion public repository (country CSV)	4468	8e304330f5a63779

S02b Microsoft AI Diffusion repo README (registry fingerprint c	3541	016d900e299594e0
S01 Microsoft Global AI Diffusion Q1 2026 report PDF	1784281	3bf9f87b63372412
S01b Microsoft Global AI Diffusion 2025 H1 report PDF (first re	7746221	a1220d061a69362b
S15 World Bank GDP current US\$ (Frictionless datasets/gdp mirr	562767	8d84ef6bcaccf740
X01 World Bank population total (Frictionless datasets/populat	552112	7d2dd6a17f5ed791
X02 Country codes / regions concordance (Frictionless datasets	134003	67b009b529330b0a

Microsoft repository vintage: git HEAD 8fe84a3b797e89846e892592dbb9d0022e9b7dd2. Permutations: 10,000 per spec (2,000 for the family max-stat), seed 20260825. Full run ledger: exploration/search_ledger.json; validation record: manifest/validation.json. Blocked sources are recorded with status blocked_by_egress in manifest/manifest.json. This note was generated by reports/generate_note.py; no numeric result was typed by hand.